

Tyler DeFoor
PA4 - Bucketsort
4/12/2017
Dr. Harris

Sequential	3
Sequential Implementation	3
Sequential Analysis	3
Sequential Graph and Table	4
Parallel	5
Parallel Implementation	5
Parallel Analysis	5
Parallel Graphs	6
Future Works	9

Sequential


Sequential Implementation

This implementation of matrix multiplication is very simple. It randomly generates numbers to fill two square matrices, then multiplies them together. The matrix multiplication portion of the code is three nested for loops. The outside goes through the rows of the result, the second goes through the columns of the result, and the final goes through the rows and columns of the left and right matrices respectively. It then has the option to print out all three matrices to be sure that everything is executing appropriately.

Sequential Analysis

Matrix multiplication with this algorithm is an $O(N^3)$ algorithm due to the nested for loops. This means that the time required to calculate the matrices grows very quickly. With 1000x1000 matrices, the execution time is 12.434 seconds. At 2700x2700 the program passes the 5 minute execution mark and continues to grow from there. The numbers that are used are too small to show the $O(N^3)$ graph, but it is consistent with what an $O(N^3)$ graph looks like in its early stages.

Sequential Graph and Table



Square Matrix Size	Execution Time	Square Matrix Size	Execution Time
1000	12.434	2000	130.444
1100	17.7073	2100	148.932
1200	24.1395	2200	176.078
1300	33.8468	2300	218.29
1400	42.2615	2400	227.498
1500	52.0867	2500	257.164
1600	65.2436	2600	296.043
1700	76.8786	2700	337.27
1800	92.5956	2800	374.922
1900	108.398	2900	432.798

Parallel

Parallel Implementation

This program implements Cannon's algorithm for parallel matrix multiplication. It requires a square matrix, a square number of processors, and for the width of the matrices to be divisible by the square root of the number of processors. This is due to the fact that the processors are formed into a grid, with each receiving its own row and column number. The matrix is then equally distributed to all of the processors.

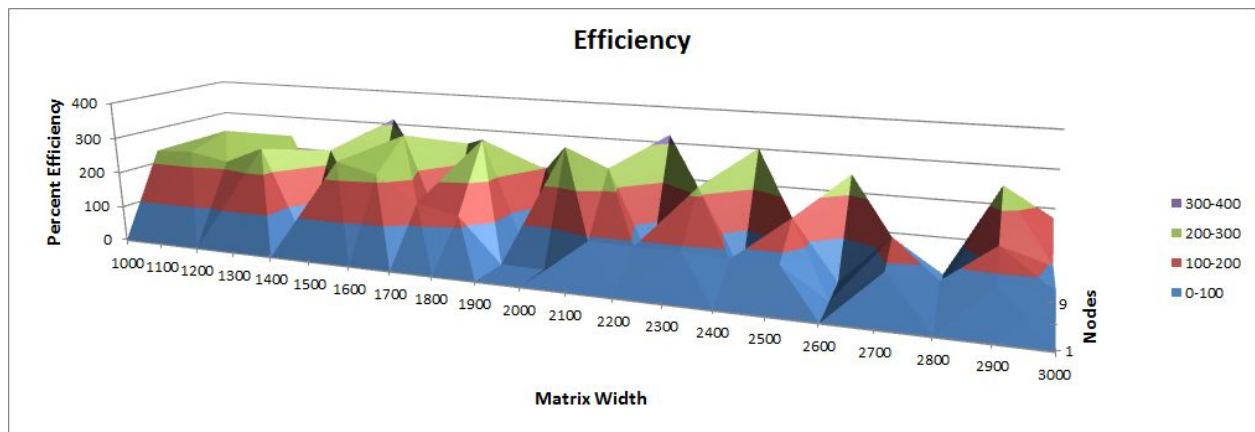
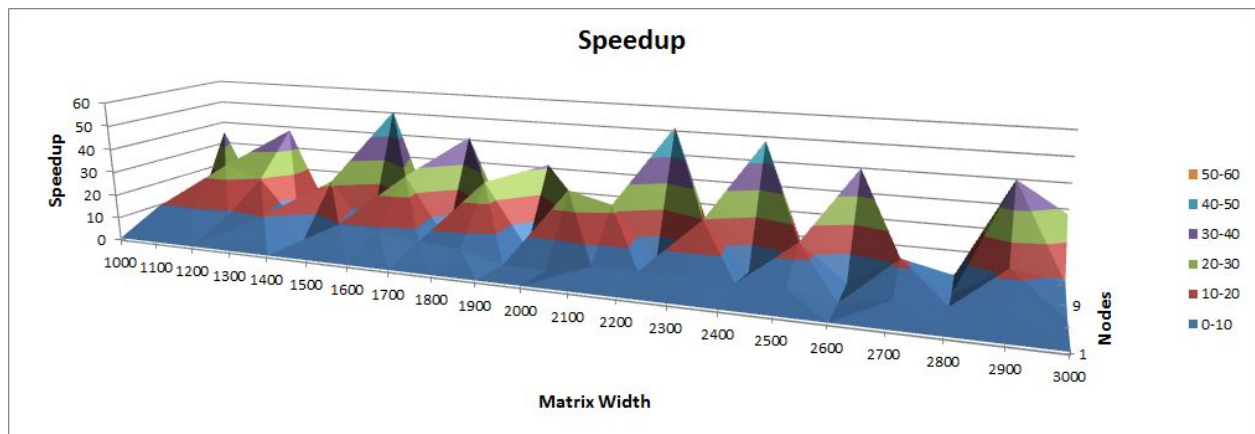
After the initial distribution, the matrices are preprocessed. All of the rows in the left matrix are shifted left by their row number, and all of columns in the right matrix are shifted up by their row number. For example, processor 2,3 will shift their left matrix's rows two to the left in the processor grid, and it will shift the right matrix's columns up three times in the processor grid. After the preprocessing, every processor's left and right matrices are multiplied together. The left matrix is passed to the processor's left neighbor, and the right matrix is passed to the processor's upper neighbor. This is repeated $N - 1$ times, where N is the square root of the number of processors.

Parallel Analysis


Matrix multiplication is an $O(N^3)$ algorithm, most of which is parallelizable. This means that the speedup and efficiency of adding cores is very large. As shown in the charts below, getting superlinear speedup was very common. It did not take long for the sequential portion to reach the five minute limit, so the timings are in comparison to a maximum of a 3,000 by 3,000 matrix. Many of the timings are not exactly the same, as the parallel could not work on the increments of 100 like the parallel did. The discrepancies are not large and each was compared against its closest matrix size.

The greatest comparative speedup came from the 2,200 by 2,200 matrix on 16 cores, with a speedup of 51.6. This was also the greatest efficiency, at 322.5023%. These superlinear speedups are due to low message size and the number of processors working on each task. Having 16 processors working on the same matrix that a single processor would is a huge time save with an $O(N^3)$ algorithm.


Parallel Graphs



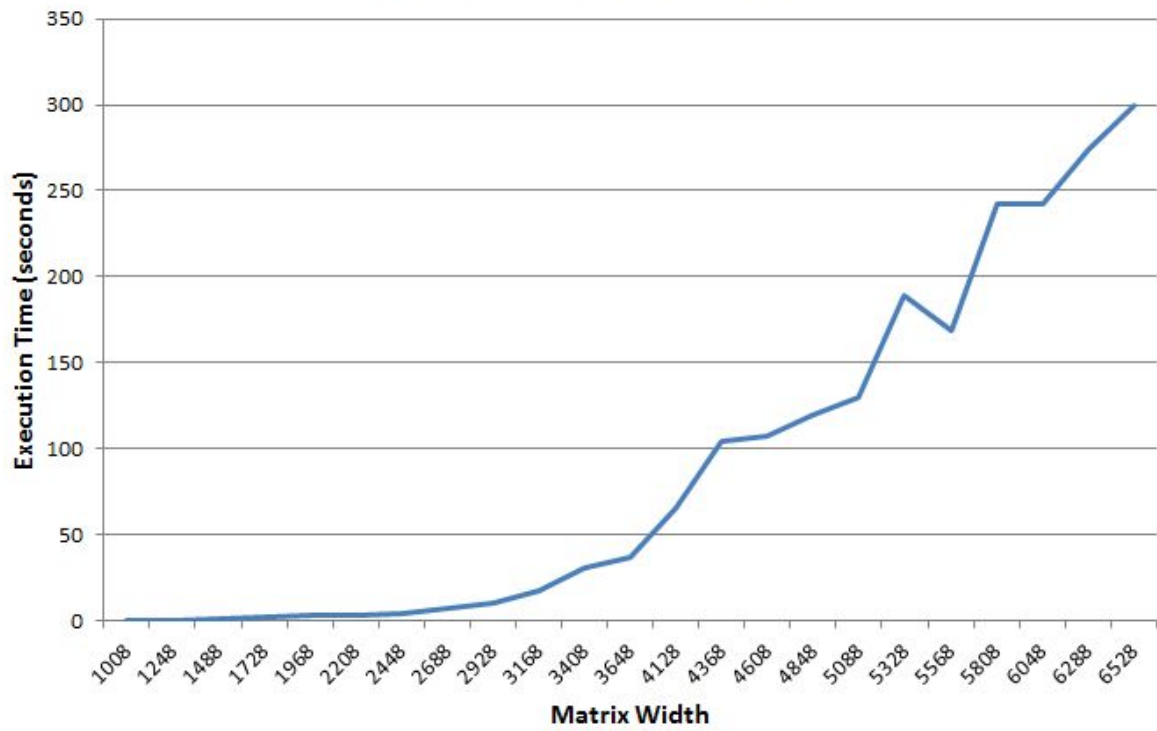
4 Node Execution Time



9 Node Execution Time



16 Node Execution Time



Future Works

I would like to revisit this in the future for two reasons. The first is to improve the program to work on matrices that are indivisible by the square processor size. The second is that the algorithm is already very scalable, so using it on clusters with more processors would be interesting to see.